

VLM Reality Check — Antigravity Prompt Engineering Pipeline

Dataset Creation Only | Stages 1–5

How to use this document Each stage is a deployable unit. Paste the prompt into Antigravity exactly as written. Do NOT proceed to the next stage until ALL test cases in the current stage pass. Treat each gate like a CI/CD merge check.

Pipeline Architecture

```
STAGE 1: Repo Scaffold & Config
    ↓ [GATE 1]
STAGE 2: Data Collection Pipeline
    ↓ [GATE 2]
STAGE 3: Vision Processing & Annotation
    ↓ [GATE 3]
STAGE 4: Adversarial Challenge Generator (9 Biases)
    ↓ [GATE 4]
STAGE 5: Multilingual Translation Pipeline
    ↓ [GATE 5]
    ✓ DATASET COMPLETE
    → data/challenges/challenges_multilingual.jsonl
    → 130K challenges × 5 languages = 650K instances
```

STAGE 1 — Repo Scaffold & Configuration

Prompt

You are a Staff ML Engineer at a top AI research lab.

Project: VLM Reality Check – a large-scale adversarial diagnostic benchmark measuring 9 systematic biases in Vision-Language Models across 130K image-pair challenges and 5 languages.

Task: Scaffold the complete project repository structure for this pipeline.
Follow production ML research conventions (similar to HuggingFace datasets repos).

Requirements:

- Monorepo layout with clear separation: `data/`, `src/`, `configs/`, `scripts/`, `outputs/`
- `configs/biases.yaml`: defines all 9 bias types with metadata:
 - `name`,
 - `difficulty_split` (easy/medium/hard summing to 1.0),
 - `ground_truth_method`,
 - `target_challenge_count`
- `Biases`: `texture`, `counting`, `spatial_relations`, `physical_plausibility`,
`temporal_reasoning`, `spurious_correlation`, `compositional_binding`,
`text_in_image`, `scale_invariance`
- `configs/pipeline.yaml`: global settings
 - `target_dataset_size: 130000`

```

compound_challenge_count: 40000
languages: [en, es, zh, hi, ar]
random_seed: 42
image_sources: [wikimedia, openimages, streetview, yfcc]
difficulty_distribution: {easy: 0.30, medium: 0.50, hard: 0.20}
- configs/languages.yaml: EN, ES, ZH, HI, AR with Helsinki-NLP model IDs
- src/__init__.py with version = "0.1.0"
- Makefile with targets: collect, process, generate, translate, validate_dataset
- requirements.txt with pinned versions for:
  torch, transformers, ultralytics, opencv-python, astropy,
  pytesseract, Pillow, datasets, scipy, pandas, tqdm,
  requests, beautifulsoup4, imagehash
- .env.example with: STREET_VIEW_API_KEY, HF_TOKEN
- README.md with ASCII pipeline diagram and quickstart

```

Return the file tree first, then generate each config file in full.
 Do not generate any src/ code yet – configs and scaffold only.

Rationale

Configs-first forces explicit decisions about taxonomy, scale, and target counts before any code is written. Errors here are cheap. Errors in Stage 4 are expensive.

GATE 1 — Test Cases

Test 1.1 — Bias taxonomy completeness

```

import yaml

with open("configs/biases.yaml") as f:
    biases = yaml.safe_load(f)

required_biases = {
    'texture', 'counting', 'spatial_relations', 'physical_plausibility',
    'temporal_reasoning', 'spurious_correlation', 'compositional_binding',
    'text_in_image', 'scale_invariance'
}
assert set(biases.keys()) == required_biases, \
    f"Missing or extra biases: {set(biases.keys()) ^ required_biases}"

for name, meta in biases.items():
    splits = meta['difficulty_split']
    assert abs(sum(splits.values()) - 1.0) < 0.01, f"{name} splits don't sum to 1.0"
    assert 'ground_truth_method' in meta, f"{name} missing ground_truth_method"
    assert 'target_challenge_count' in meta, f"{name} missing target_challenge_count"

total = sum(b['target_challenge_count'] for b in biases.values())
assert 85000 <= total <= 95000, f"Single-bias targets should sum ~90K, got {total}"
print("GATE 1.1 PASSED")

```

PASS: Prints "GATE 1.1 PASSED"

FAIL: Any assert fails → fix the offending entry in configs/biases.yaml only

Test 1.2 — Pipeline config has all required keys

```

import yaml

with open("configs/pipeline.yaml") as f:
    cfg = yaml.safe_load(f)

required_keys = {
    'target_dataset_size', 'compound_challenge_count', 'languages',
    'random_seed', 'image_sources', 'difficulty_distribution'
}

```

```
missing = required_keys - set(cfg.keys())
assert not missing, f"Missing pipeline config keys: {missing}"
assert cfg['target_dataset_size'] == 130000
assert cfg['compound_challenge_count'] == 40000
assert set(cfg['languages']) == {'en', 'es', 'zh', 'hi', 'ar'}
assert cfg['random_seed'] == 42
print("GATE 1.2 PASSED")

PASS: Prints "GATE 1.2 PASSED"
FAIL: Missing key → add to configs/pipeline.yaml
```

Test 1.3 — Makefile targets resolve

```
make -n collect && make -n process && make -n generate && \
make -n translate && make -n validate_dataset
echo "GATE 1.3 PASSED"

PASS: All targets resolve without "No rule to make target" error
FAIL: Missing target → add to Makefile
```

STAGE 2 — Data Collection Pipeline

Prompt

You are a Staff ML Engineer. Stage 1 scaffold is complete.

Task: Build src/collection/ – the image scraping and download pipeline targeting 500K images from 4 sources.

Architecture: 4 source collectors, 1 unified pipeline with deduplication.

Implement:

src/collection/__init__.py

src/collection/base_collector.py
- Abstract class BaseCollector
- ImageRecord = dataclass(image_id, url, local_path, source, metadata: dict, license)
- Abstract method: fetch_batch(n: int) → List[ImageRecord]
- Built-in retry logic: 3 retries, exponential backoff (1s, 2s, 4s)
- Rate limiting: configurable requests_per_second (default 2)
- Deduplication: perceptual hash (pHash) stored in SQLite at data/raw/seen_hashes.db
Threshold: hamming distance < 10 = duplicate, skip it

src/collection/wikimedia_collector.py
- Wikimedia Commons API – free, no key required
- Filter: CC-licensed only, min resolution 512x512
- Categories: ["outdoor scenes", "urban streets", "animals", "everyday objects"]
- Returns ImageRecord with license field populated

src/collection/openimages_collector.py
- Open Images v7 CSV manifest (assume at data/raw/yfcc_metadata.csv)
- Filter: has_bbox=True, min 50 images per object class
- Returns stratified sample across object categories

src/collection/streetview_collector.py
- Google Street View Static API (key from .env STREET_VIEW_API_KEY)
- Global city grid: 100 major cities, 50 GPS points per city

- Per location: 4 headings (0° , 90° , 180° , 270°) for spatial pair challenges
- Graceful fallback: if API key missing → log warning and skip this source entirely

```
src/collection/yfcc_collector.py
- YFCC100M metadata CSV at data/raw/yfcc_metadata.csv
- Filter: has_gps=True, is_outdoor=True, CC license, year < 2020
- Metadata dict must include: timestamp, lat, lon, camera_model
```

```
src/collection/pipeline.py
- CollectionPipeline class:
  - Runs all 4 collectors in parallel via ThreadPoolExecutor (max_workers=4)
  - Deduplicates across sources using pHash
  - Saves images to data/raw/images/{source}/{image_id}.jpg
  - Saves manifest to data/raw/manifest.jsonl
    Each line: {image_id}, source, local_path, url, license, metadata
  - Checkpoint-resume: reads already-completed image_ids from manifest.jsonl
    on startup, skips re-downloading those IDs
  - Progress: tqdm bar showing total collected, per-source counts, duplicates skipped
  - Target: 500K images total
```

Use type hints everywhere. Docstrings on all public methods.

Rationale

pHash deduplication prevents the same image appearing as both image A and image B in a challenge pair, which would corrupt ground truth. Checkpoint-resume is mandatory at 500K scale — collection will not finish in one uninterrupted session.

GATE 2 — Test Cases

Test 2.1 — ImageRecord schema is correct

```
from src.collection.base_collector import ImageRecord
import dataclasses

fields = {f.name for f in dataclasses.fields(ImageRecord)}
required = {'image_id', 'url', 'local_path', 'source', 'metadata', 'license'}
assert required.issubset(fields), f"Missing fields: {required - fields}"
print("GATE 2.1 PASSED")

PASS: Prints "GATE 2.1 PASSED"
FAIL: Missing field → add to ImageRecord dataclass in base_collector.py
```

Test 2.2 — pHash deduplication catches near-identical images

```
import imagehash
from PIL import Image
import numpy as np

# Exact duplicate
img = Image.fromarray(np.random.randint(0, 255, (512, 512, 3), dtype=np.uint8))
h1 = imagehash.phash(img)
h2 = imagehash.phash(img)
assert (h1 - h2) < 10, "Exact duplicate not caught by pHash threshold"

# Clearly different image
img2 = Image.fromarray(np.random.randint(0, 255, (512, 512, 3), dtype=np.uint8))
h3 = imagehash.phash(img2)
# Different images should typically have distance > 10
# (probabilistic but will pass with overwhelming probability for random images)
print("GATE 2.2 PASSED")

PASS: No assertion error
FAIL: imagehash not installed → add to requirements.txt and reinstall
```

Test 2.3 — Checkpoint-resume skips already-downloaded images

```
from src.collection.pipeline import CollectionPipeline
from unittest.mock import patch

pipeline = CollectionPipeline(config_path="configs/pipeline.yaml")

# Pre-write 3 image IDs to manifest as if already downloaded
import json
with open("data/raw/manifest.jsonl", "w") as f:
    for i in range(3):
        f.write(json.dumps({"image_id": f"img_{i:04d}", "source": "wikimedia",
                            "local_path": f"data/raw/images/img_{i:04d}.jpg",
                            "url": "http://example.com", "license": "CC", "metadata": {}}) + "\n")

with patch.object(pipeline, '_download_image', return_value=True) as mock_dl:
    pipeline.run(target_n=3, resume=True)
    assert mock_dl.call_count == 0, f"Expected 0 downloads, got {mock_dl.call_count}"

print("GATE 2.3 PASSED")

PASS: 0 redundant downloads triggered
FAIL: Re-downloading → fix checkpoint loading logic in pipeline.py __init__
```

STAGE 3 — Vision Processing & Annotation Pipeline

Prompt

You are a Staff ML Engineer and Computer Vision researcher.

Stage 2 is complete. Images are at data/raw/images/, manifest at data/raw/manifest.jsonl.

Task: Build src/processing/ – the automated annotation pipeline that extracts structured metadata from each image. This metadata is the ground truth source for all 9 bias challenge generators in Stage 4.

Critical: each processor must be independently importable and runnable.

Use result caching (SQLite) so re-runs don't reprocess existing images.

Implement:

src/processing/__init__.py

src/processing/base_processor.py

- Abstract ProcessorBase:
 - Abstract method: process(image_path: str) -> dict
 - Cache layer: SQLite at data/processed/annotation_cache.db
 - Key: image_id + processor_name. Returns cached dict if exists.
 - Error handling: on exception, return {"error": str(e), "skipped": True}
 - Never crash the pipeline on a single bad image.

src/processing/object_detector.py ← Used by: counting, spurious, scale

- YOLOv8x model via ultralytics
- Returns: {detections: List[{"label", "bbox", "confidence", "area_fraction"}]}
- Only include detections with confidence > 0.7
- area_fraction = bbox_area / (image_width * image_height)
- Batch inference: process images in batches of 32

src/processing/depth_estimator.py ← Used by: spatial_relations, scale

- MiDaS DPT-Large via transformers
- Returns: {left_depth_mean, right_depth_mean, top_depth_mean, bottom_depth_mean} (average depth in each image quadrant)
- Normalize depth map to [0, 1] before computing quadrant means

src/processing/shadow_detector.py ← Used by: physical_plausibility

- Two steps:
 - Step A – Shadow detection:
 - OpenCV HSV thresholding ($S < 50$, $V < 80$) + contour analysis
 - Returns: shadow_angle_detected (degrees from vertical, None if no shadow found)
 - Step B – Expected shadow angle:
 - Astropy: given EXIF GPS (lat, lon) + EXIF DateTimeOriginal → compute sun azimuth and elevation via AltAz frame
 - Shadow angle = sun_azimuth + 180° (shadow opposite sun)
 - Returns: shadow_angle_expected (degrees), sun_elevation (degrees)
 - Step C – Plausibility check:
 - angle_delta = |shadow_angle_detected - shadow_angle_expected|
 - is_physically_plausible = (angle_delta < 25°) if both angles available else None
- Fallback: if EXIF GPS or timestamp missing → return all fields as None, skipable=True

src/processing/ocr_extractor.py ← Used by: text_in_image

- Tesseract via pytesseract
- Returns: {
 - has_text: bool,
 - text_blocks: List[{text, bbox, confidence}],
 - text_is_location_relevant: bool ← True if any block matches street/road/place patterns
}
- Filter: only blocks with confidence > 60 and len(text.strip()) > 2

src/processing/texture_analyzer.py ← Used by: texture_bias

- Returns: {
 - edge_density: float, ← Canny edge pixel fraction
 - dominant_texture_freq: float, ← Peak Gabor filter response frequency
 - silhouette_extractable: bool ← True if edge_density > 0.05 and largest contour area > 10% of image
}

src/processing/temporal_extractor.py ← Used by: temporal_reasoning

- Parse EXIF DateTimeOriginal
- Returns: {
 - timestamp: str (ISO format, None if missing),
 - time_of_day: str (dawn/morning/noon/afternoon/dusk/night),
 - season: str (spring/summer/autumn/winter),
 - construction_score: float ← fraction of YOLO detections that are machinery/scaffolding/construction labels
}

src/processing/annotation_pipeline.py

- AnnotationPipeline: runs all 7 processors on each image, merges dicts
- Output: data/processed/annotations.jsonl
 - Each line: {image_id, detections, depth, shadow, ocr, texture, temporal}
- Skip if image_id already annotated (cache check)
- GPU batching for YOLO and MiDaS (batch_size=32)
- tqdm progress bar with ETA
- On completion: print summary – total processed, skipped (cached), errored

Rationale

The shadow detector (Astropy + OpenCV) is the key innovation that makes physical plausibility ground truth scientifically automated rather than manually labeled. The skipable=True fallback is critical — images without EXIF GPS must be cleanly excluded from physical_plausibility challenges, not cause pipeline crashes.

GATE 3 — Test Cases

Test 3.1 — Shadow plausibility is deterministically correct

```

from src.processing.shadow_detector import ShadowDetector

detector = ShadowDetector()
# Known case: NYC (40.71°N, 74.01°W), summer solstice noon
# Sun is due south, elevation ~72° → expected shadow points north (~180° azimuth)
result = detector.compute_expected_shadow(
    lat=40.7128, lon=-74.0060,
    timestamp="2024-06-21T12:00:00"
)
assert result['shadow_angle_expected'] is not None
assert 160 < result['shadow_angle_expected'] < 200, \
    f"Expected shadow ~180° (north), got {result['shadow_angle_expected']}"
assert result['sun_elevation'] > 60, "Sun should be high at NYC noon in June"
print("GATE 3.1 PASSED")

PASS: Prints "GATE 3.1 PASSED"
FAIL: Wrong angle → check Astropy AltAz coordinate frame and timezone handling

```

Test 3.2 — Object detector returns correct schema

```

from src.processing.object_detector import ObjectDetector
import os

detector = ObjectDetector()
# Use any real image from data/raw/images/ or a test fixture
test_img = "tests/fixtures/sample.jpg" # place any outdoor JPEG here
result = detector.process(test_img)

assert 'detections' in result
assert isinstance(result['detections'], list)
for det in result['detections']:
    assert 'label' in det and 'confidence' in det and 'bbox' in det and 'area_fraction' in det
    assert det['confidence'] > 0.7, "Low-confidence detections should be filtered"
    assert 0 < det['area_fraction'] <= 1.0
print("GATE 3.2 PASSED")

PASS: Prints "GATE 3.2 PASSED"
FAIL: Schema mismatch → fix return dict structure in object_detector.py

```

Test 3.3 — Annotation pipeline output is complete and cacheable

```

import json, os
from src.processing.annotation_pipeline import AnnotationPipeline

# Run on 5 fixture images
pipeline = AnnotationPipeline()
pipeline.run(image_ids=["test_001", "test_002", "test_003", "test_004", "test_005"],
            images_dir="tests/fixtures/")

required_keys = {'image_id', 'detections', 'depth', 'shadow', 'ocr', 'texture', 'temporal'}
with open("data/processed/annotations.jsonl") as f:
    records = [json.loads(l) for l in f]

assert len(records) == 5
for rec in records:
    missing = required_keys - set(rec.keys())
    assert not missing, f"Missing keys {missing} in record {rec.get('image_id')}"

# Run again — should use cache, not reprocess
import time
t0 = time.time()
pipeline.run(image_ids=["test_001", "test_002", "test_003", "test_004", "test_005"],
            images_dir="tests/fixtures/")

```

```

    images_dir="tests/fixtures/")
elapsed = time.time() - t0
assert elapsed < 2.0, f"Cache not working - reprocessing took {elapsed:.1f}s"
print("GATE 3.3 PASSED")

PASS: Prints "GATE 3.3 PASSED"
FAIL: Missing keys → add null fallback in annotation_pipeline.py merge step
      Cache not working → fix cache key lookup in base_processor.py

```

STAGE 4 — Adversarial Challenge Generator (9 Biases)

Prompt

You are a Staff ML Research Engineer.
Stages 1–3 complete. Annotations at data/processed/annotations.jsonl.

Task: Build src/generators/ – the adversarial challenge generation engine.
This is the scientific core of the benchmark.

CRITICAL CONSTRAINT: Each challenge must vary EXACTLY ONE factor between Image A and Image B. Challenges that vary multiple factors are confounded and scientifically invalid. Reject them – do not patch them.

Implement:

```

src/generators/base_generator.py
- Challenge = dataclass(
    challenge_id: str,
    bias_type: str,
    difficulty: str,          # easy / medium / hard
    image_a_id: str,
    image_b_id: str,          # same as image_a_id for single-image transforms
    question_template: str,    # uses {variable} slots
    correct_answer: str,
    distractor_answers: List[str],
    ground_truth_method: str, # e.g. "yolo_count", "astropy_shadow", "ocr_text"
    confound_check_passed: bool,
    metadata: dict
)
- Abstract ChallengeGenerator:
- Abstract: generate_challenge(annotations: List[dict]) -> Challenge | None
- validate_single_factor_isolation(challenge: Challenge) -> bool
  Checks: only the bias-relevant attribute differs between images.
  Returns False → generator must return None (rejected).

```

```

src/generators/texture_generator.py
- Pairs: real photo of object vs OpenCV-generated silhouette of the SAME image
  (Canny edges + filled contour = silhouette, same object, texture removed)
- Easy: high-edge-density object (dog, car) – shape is clear
- Hard: low-edge-density, ambiguous shape (cat vs dog silhouette)
- Question: "Do both images show the same type of object? Answer Yes or No."
- Correct answer: always Yes (same object, texture removed)
- Bias signal: model answers No when texture cues removed

```

```

src/generators/counting_generator.py
- Pairs two images with YOLO-verified counts N vs M, same object category
- Require |N - M| >= 2
- Easy: |N - M| >= 4

```

- Medium: $|N - M| = 2$ or 3
- Hard: N=5, M=6 (both above subitizing limit of 4)
- Question: "Which image has more {object_category}? Answer A or B."
- Correct answer: deterministic from stored YOLO counts

`src/generators/spatial_generator.py`

- Takes one image, creates adversarial version via:
 - Horizontal flip → left/right inverted
 - Vertical flip → above/below inverted
- Question: "Are these the same scene from the same viewpoint? Answer Yes or No."
- Correct answer: always No
- Confound check: verify only spatial orientation changed (same pixel content, flipped)

`src/generators/physics_generator.py`

- Source images: only those where shadow.is_physically_plausible = True and shadow.shadow_angle_detected is not None
- Creates adversarial copy: rotate detected shadow region by 120-150° using OpenCV affine transform (shadow pixels only, not full image)
- Easy: shadow points directly opposite correct direction (180° off)
- Medium: shadow 60° off expected angle
- Hard: shadow 30° off (subtly wrong, still plausible-looking)
- Original image: Question "Is the lighting physically possible? Yes or No." → Yes
- Modified image: same question → No
- Confound check: only shadow angle changed, no other pixel regions modified

`src/generators/temporal_generator.py`

- Groups 3 images by temporal sequence:
 - Option A: dawn → noon → dusk (time_of_day progression, same scene type)
 - Option B: construction_score low → medium → high (same GPS cluster)
- Question: "Order these images chronologically. Answer as A,B,C or A,C,B etc."
- Correct answer: the correct temporal ordering string
- Hard: 3 images all within 2 hours (subtle lighting change)

`src/generators/spurious_generator.py`

- Pairs image of entity in TYPICAL context vs ATYPICAL context
 - Example: cow on grass (typical) vs cow on beach (atypical)
- Same YOLO label, different background context (determined from scene metadata)
- Question: "What is the main {category} in this image? Name it."
- Correct answer: same object label for BOTH images
- Bias signal: model gives different answers based on background

`src/generators/compositional_generator.py`

- Selects images with 2+ YOLO-detected objects each having distinct dominant colors (computed via color histogram per bbox)
- Generates correct description and 3 foil descriptions (swapped attribute bindings)
- Question: "Which description correctly matches the image?"
- Correct answer: the non-swapped description
- Hard: objects with similar colors (dark blue vs dark green)

`src/generators/text_image_generator.py`

- Pairs two images where ocr.text_is_location_relevant = True
- OCR-extracted text differs between image A and image B
- Question: "Are these the same location? Use only the visible text to decide.
Answer Yes or No."
- Correct answer: No (different text = different location)
- Confound check: images must be visually similar except for text content

`src/generators/scale_generator.py`

- Takes one image with a clearly detected object (area_fraction 0.1-0.4)
- Creates zoomed version: crop to 1/zoom_factor of image centered on object bbox
 - Zoom factor: randomly chosen from [5x, 8x, 10x, 15x]
- Question: "Do these show the same type of object? Answer Yes or No." → Yes
- Second question: "Are these objects the same real-world size? Answer Yes or No." → No
- This split isolates recognition bias from scale bias

```

src/generators/compound_generator.py
- Combines 2 single-bias generators sequentially on same image pair
- Supported combinations (must be visually coherent):
  spatial + physics, counting + spurious, texture + scale, compositional + text
- Records: bias_types_combined: ["spatial_relations", "physical_plausibility"]
- difficulty: always "hard"
- metadata: includes expected_failure_mode per bias component

src/generators/pipeline.py
- GenerationPipeline:
  - Loads annotations.jsonl
  - Runs each generator to hit per-bias target counts from configs/biases.yaml
  - Compound generator runs last (depends on single-bias images being validated)
  - Total target: 90K single-bias + 40K compound = 130K challenges
  - Saves to data/challenges/challenges.jsonl
  - Rejects any challenge where confound_check_passed = False (log rejection)
  - On completion prints: generated per bias, rejected per bias, total

```

Rationale

The `confound_check_passed` rejection (not patching) is the scientific boundary condition of this benchmark. The compound generator must run last because compound challenges combine already-validated single-bias image pairs — using unvalidated pairs would silently contaminate the compound set. Scale invariance is split into two sub-questions deliberately — this is the only way to separate object recognition from size estimation.

GATE 4 — Test Cases

Test 4.1 — Confounded challenge is rejected, not accepted

```

from src.generators.spatial_generator import SpatialGenerator

gen = SpatialGenerator()
# Mock annotation pair where BOTH spatial orientation AND object count differ
# (would be confounded — testing two things at once)
confounded_annotations = [
    {"image_id": "img_a", "detections": {"detections": [{"label": "cat"}]*3},
     "depth": {"left_depth_mean": 0.3}},
    {"image_id": "img_b", "detections": {"detections": [{"label": "cat"}]*7}, # count also differs
     "depth": {"left_depth_mean": 0.7}},
]
# Patch generate to use these annotations, then check confound check fails
result = gen.generate_challenge(confounded_annotations)
if result is not None:
    assert result.confound_check_passed == False,
        "Confounded challenge must have confound_check_passed=False"
print("GATE 4.1 PASSED")

PASS: result is None OR result.confound_check_passed == False
FAIL: Confounded challenge accepted as valid → fix validate_single_factor_isolation()

```

Test 4.2 — Counting generator assigns difficulty and answer correctly

```

from src.generators.counting_generator import CountingGenerator

gen = CountingGenerator()
# 3 cats vs 7 cats → easy (|diff|=4), answer=B
ann_3 = {"image_id": "img_a", "detections": {"detections": [{"label": "cat", "confidence": 0.9}]*3}}
ann_7 = {"image_id": "img_b", "detections": {"detections": [{"label": "cat", "confidence": 0.9}]*7}}
challenge = gen.generate_challenge([ann_3, ann_7])
assert challenge is not None
assert challenge.correct_answer == "B", "B has more cats"
assert challenge.difficulty == "easy", "|3-7|=4 should be easy"

# 5 cats vs 6 cats → hard (both >4, subitizing zone)

```

```

ann_5 = {"image_id": "img_c", "detections": {"detections": [{"label": "cat", "confidence": 0.9}] * 5}}
ann_6 = {"image_id": "img_d", "detections": {"detections": [{"label": "cat", "confidence": 0.9}] * 6}}
challenge_hard = gen.generate_challenge([ann_5, ann_6])
assert challenge_hard.difficulty == "hard", "|5-6|=1 and both >4 should be hard"
print("GATE 4.2 PASSED")

```

PASS: Prints "GATE 4.2 PASSED"

FAIL: Wrong answer or difficulty → fix ground truth and difficulty assignment logic

Test 4.3 — Final dataset distribution matches target

```

import json
from collections import Counter

challenges = [json.loads(l) for l in open("data/challenges/challenges.jsonl")]
total = len(challenges)
assert total >= 125_000, f"Only {total} challenges generated (target 130K)"

bias_counts = Counter(c['bias_type'] for c in challenges)
difficulty_counts = Counter(c['difficulty'] for c in challenges)

# No single bias > 15% of total
for bias, count in bias_counts.items():
    frac = count / total
    assert frac <= 0.15, f"{bias} is {frac:.1%} of dataset - too dominant"

# Difficulty distribution within ±5% of target
easy_frac = difficulty_counts['easy'] / total
medium_frac = difficulty_counts['medium'] / total
hard_frac = difficulty_counts['hard'] / total
assert 0.25 <= easy_frac <= 0.35, f"Easy {easy_frac:.2%} outside [25%, 35%]"
assert 0.45 <= medium_frac <= 0.55, f"Medium {medium_frac:.2%} outside [45%, 55%]"
assert 0.15 <= hard_frac <= 0.25, f"Hard {hard_frac:.2%} outside [15%, 25%]"

# Confirm compound challenges exist
compound = [c for c in challenges if c['bias_type'] == 'compound']
assert len(compound) >= 35_000, f"Only {len(compound)} compound challenges (target 40K)"
print("GATE 4.3 PASSED")

```

PASS: Prints "GATE 4.3 PASSED"

FAIL: Distribution off → adjust target_challenge_count in configs/biases.yaml
and re-run generation pipeline

STAGE 5 — Multilingual Translation Pipeline

Prompt

You are a Staff NLP Engineer.

Stage 4 complete. Challenges at data/challenges/challenges.jsonl.

Task: Build src/translation/ – translate all 130K challenge questions into 4 additional languages (ES, ZH, HI, AR) using Helsinki-NLP models locally.

Critical design: template-first translation, NOT sentence-by-sentence.

Translate the question TEMPLATE (with {variable} slots preserved), then slot-fill variables after. This prevents translation drift and variable corruption.

Implement:

```
src/translation/__init__.py
```

```
src/translation/question_templates.py
- Define 3-5 question templates PER BIAS TYPE in English
- Variables in {} are preserved through translation unchanged
- Example for counting:
    TEMPLATES['counting'] = [
        "Which image has more {object_category}? Answer A or B.",
        "Image A shows {count_a} {object_category}. Image B shows {count_b}.
        Which has more? Answer A or B.",
    ]
- Templates for all 9 biases + compound. Total: ~50 templates.
- Each template must have: template_id, bias_type, variables: List[str],
  valid_answers: List[str]
```

```
src/translation/helsinki_translator.py
- Loads Helsinki-NLP/opus-mt-{src}-{tgt} models from HuggingFace:
  EN→ES: Helsinki-NLP/opus-mt-en-es
  EN→ZH: Helsinki-NLP/opus-mt-en-zh
  EN→HI: Helsinki-NLP/opus-mt-en-hi
  EN→AR: Helsinki-NLP/opus-mt-en-ar
- Template translation method:
  1. Replace {variable_name} with [SLOT_0], [SLOT_1], etc.
  2. Translate the de-slotted template string
  3. Restore [SLOT_N] back to {variable_name}
- Batch size: 32 templates per forward pass
- Cache: if (template_id, target_lang) already in SQLite cache → return cached
- Returns: translated_template string with {} slots intact
```

```
src/translation/script_validator.py
- Validates translated text has correct script characters:
  ES: basic Latin – assert no [SLOT] tokens remain
  ZH: assert contains CJK unicode range (U+4E00–U+9FFF)
  HI: assert contains Devanagari range (U+0900–U+097F)
  AR: assert contains Arabic range (U+0600–U+06FF)
- For AR (RTL): prepend Unicode RLM marker \u200F for correct rendering
- For ZH: strip extra spaces between characters
```

```
src/translation/translation_pipeline.py
- Loads challenges.jsonl
- For each challenge:
  1. Select appropriate template by bias_type
  2. Translate template for all 4 non-English languages (cached)
  3. Slot-fill {variables} from challenge metadata
  4. Validate script per language via script_validator
- Output: data/challenges/challenges_multilingual.jsonl
  Each record is the original challenge dict + new field:
  "questions": {"en": "...", "es": "...", "zh": "...", "hi": "...", "ar": "..."}
- Log: cache hit rate, per-language template count, any failed validations
```

```
src/translation/human_validation_sampler.py
- Randomly samples (seed=42) 100 challenges per language = 500 total
- Output: data/validation/human_validation_sample.csv
  Columns: challenge_id, bias_type, language, en_question, translated_question,
            is_fluent, is_accurate, notes
            (is_fluent / is_accurate / notes left blank for human annotators to fill)
```

Rationale

Template-first translation is mandatory here. Free-form translation of complete sentences like "Which image has more cats?" will sometimes translate "cats" to a localized equivalent, breaking the variable slot that should have been "{object_category}". The slot-masking approach guarantees object names, answer choices (A/B), and numbers survive translation intact. Arabic RTL requires the unicode RLM marker or the question renders garbled in most interfaces.

GATE 5 — Test Cases

Test 5.1 — Variable slots survive all 4 translations

```
from src.translation.helsinki_translator import HelsinkiTranslator

translator = HelsinkiTranslator()
template = "Which image has more {object_category}? Answer A or B."

for lang in ['es', 'zh', 'hi', 'ar']:
    result = translator.translate_template(template, target_lang=lang)
    assert '{object_category}' in result, \
        f"Slot '{object_category}' lost in {lang} translation: '{result}'"
    assert '[SLOT' not in result, \
        f"SLOT token not restored in {lang}: '{result}'"

print("GATE 5.1 PASSED")
```

PASS: All 4 languages preserve {object_category} and restore SLOT tokens
FAIL: Slot eaten → fix masking/restoring logic in helsinki_translator.py

Test 5.2 — Script validation catches wrong-script outputs

```
from src.translation.script_validator import ScriptValidator

validator = ScriptValidator()

# Correct scripts should pass
assert validator.validate("Hola mundo", lang="es") == True
assert validator.validate("你好世界", lang="zh") == True
assert validator.validate("नमस्ते दुनिया", lang="hi") == True
assert validator.validate("مرحبا بالعالم", lang="ar") == True

# Wrong script should fail
assert validator.validate("Hello world", lang="zh") == False, \
    "English text should fail Chinese validation"
assert validator.validate("Hello world", lang="ar") == False, \
    "English text should fail Arabic validation"

print("GATE 5.2 PASSED")
```

PASS: Prints "GATE 5.2 PASSED"
FAIL: Validation too loose → tighten unicode range checks in script_validator.py

Test 5.3 — Final multilingual output is complete

```
import json

records = [json.loads(l) for l in open("data/challenges/challenges_multilingual.jsonl")]
langs = ['en', 'es', 'zh', 'hi', 'ar']

issues = []
for rec in records[:2000]: # spot-check first 2K
    for lang in langs:
        q = rec.get('questions', {}).get(lang, '')
        if not q or len(q.strip()) < 5:
            issues.append(f"challenge {rec['challenge_id']} lang {lang} is empty/short")

assert len(issues) == 0, f"{len(issues)} translation issues:\n" + "\n".join(issues[:5])

# Verify total record count
assert len(records) >= 125_000, f"Only {len(records)} records (expected 130K)"
print(f"GATE 5.3 PASSED - {len(records)} multilingual challenges generated")
```

```
PASS: Prints "GATE 5.3 PASSED - X multilingual challenges generated"
FAIL: Empty translations → check which bias_type × language combination is failing
      and verify template coverage in question_templates.py
```

✓ Dataset Complete — Final Smoke Test

Run this after all 5 gates pass

```
"""
Dataset completeness smoke test.
Verifies the final deliverable before handing off to evaluation.
"""

import json
from collections import Counter
from pathlib import Path

print("Running final dataset smoke test...")

# 1. Check all files exist
assert Path("data/raw/manifest.jsonl").exists(), "manifest.jsonl missing"
assert Path("data/processed/annotations.jsonl").exists(), "annotations.jsonl missing"
assert Path("data/challenges/challenges.jsonl").exists(), "challenges.jsonl missing"
assert Path("data/challenges/challenges_multilingual.jsonl").exists(), "multilingual missing"
assert Path("data/validation/human_validation_sample.csv").exists(), "validation sample missing"

# 2. Check image count
manifest = [json.loads(l) for l in open("data/raw/manifest.jsonl")]
assert len(manifest) >= 400_000, f"Only {len(manifest)} images collected (target 500K)"

# 3. Check challenge count and structure
challenges = [json.loads(l) for l in open("data/challenges/challenges_multilingual.jsonl")]
assert len(challenges) >= 125_000, f"Only {len(challenges)} challenges (target 130K)"

required_challenge_keys = {
    'challenge_id', 'bias_type', 'difficulty',
    'image_a_id', 'image_b_id', 'correct_answer',
    'ground_truth_method', 'confound_check_passed', 'questions'
}
for c in challenges[:100]: # spot-check
    missing = required_challenge_keys - set(c.keys())
    assert not missing, f"Challenge {c['challenge_id']} missing keys: {missing}"

# 4. Check language coverage
for c in challenges[:100]:
    for lang in ['en', 'es', 'zh', 'hi', 'ar']:
        q = c['questions'].get(lang, '')
        assert q and len(q) > 4, f"Missing {lang} in challenge {c['challenge_id']}"

# 5. Check no confounded challenges slipped through
confounded = [c for c in challenges if not c.get('confound_check_passed', True)]
assert len(confounded) == 0, f"{len(confounded)} confounded challenges in dataset!"

# 6. Summary
bias_dist = Counter(c['bias_type'] for c in challenges)
diff_dist = Counter(c['difficulty'] for c in challenges)
print("\n" + "="*60)
print("DATASET SMOKE TEST PASSED")
print("="*60)
```

```

print(f"Images collected: {len(manifest):>10,}")
print(f"Total challenges: {len(challenges):>10,}")
print(f"Languages per challenge: 5 (EN, ES, ZH, HI, AR)")
print(f"Total question instances: {len(challenges)*5:>8,}")
print(f"\nBias distribution:")
for bias, count in sorted(bias_dist.items()):
    print(f" {bias}<28} {count:>7,} ({count/len(challenges):.1%})")
print(f"\nDifficulty distribution:")
for diff, count in sorted(diff_dist.items()):
    print(f" {diff}<10} {count:>7,} ({count/len(challenges):.1%})")
print("\nDataset ready for evaluation pipeline.")

```

Summary

Stage	Deliverable	Key File
1	Repo + configs	configs/biases.yaml, configs/pipeline.yaml
2	500K images	data/raw/manifest.jsonl
3	Vision annotations	data/processed/annotations.jsonl
4	130K challenges	data/challenges/challenges.jsonl
5	650K multilingual instances	data/challenges/challenges_multilingual.jsonl

5 stages · 15 test gates · 1 integration test